

CCAF



Pass

93.33% (56/60)

DOMAIN BREAKDOWN

code_exploration	13/15 (87%)
customer_support	14/15 (93%)
extraction_pipeline	15/15 (100%)
research_pipeline	14/15 (93%)

Share

Open saved result

Missed (4)

All (60)

code_exploration

Your codebase exploration tool stores session IDs to allow engineers to continue investigations across work sessions. An engineer spent an hour yesterday analyzing a legacy authentication module, building context about its architecture and dependencies. They want to continue today. The session ID is valid, but version control shows 3 of the 12 files the agent previously read were modified overnight by a teammate's merge.

What approach best balances efficiency and accuracy?

Your answer: B. Start a fresh session to ensure the agent works with current codebase state without stale assumptions

Correct: C. Resume the session and inform the agent which specific files changed for targeted re-analysis

code_exploration

An engineer's exploration subagent spent 30 minutes analyzing a legacy payment system, reading 47 files and documenting data flows. The session was interrupted when the engineer's connection dropped. While away, a teammate merged a PR that renamed two utility functions. The engineer wants to continue the same exploration.

What's the most effective approach?

Your answer: C. Launch a fresh subagent with a summary of prior findings.

Correct: D. Resume the subagent from its previous transcript and inform it about the renamed functions.

customer_support

A customer raises three separate issues during one session: a refund inquiry (turns 1-15), a subscription question (turns 16-30), and a payment method update (turns 31-45). At turn 48, the customer asks "What happened with my refund?" The conversation is approaching context limits.

What strategy best maintains the agent's ability to address all issues throughout the session?

Your answer: A. Extract and persist structured issue data (order IDs, amounts, statuses) into a separate context layer.

Correct: C. Summarize earlier turns into a narrative description, preserving full message history only for the active issue.

research_pipeline

Production reviews reveal inconsistent handling of uncertainty in final reports. Sometimes conflicting subagent findings are synthesized into a single confident statement (losing nuance), while other times reports over-hedge with excessive qualifications (becoming unhelpful). When the web search agent returns "industry analysts estimate \$50B market size (methodology varies)" and the document analysis agent returns "peer-reviewed study estimates 35B(±7B, 95% CI)," the coordinator either picks one arbitrarily or produces vague statements like "the market may be 35B-50B depending on factors."

What systematic approach best addresses this?

Your answer: B. Implement a confidence calibration layer that normalizes subagent uncertainty expressions to standardized probability scores (0.0-1.0), then weight-average findings by their calibrated confidence.

Correct: C. Instruct the synthesis agent to structure reports with explicit sections distinguishing well-established findings from contested ones, preserving original source characterizations and methodological context.

Save your progress from here

Email + a 6-digit PIN attaches future sittings to an account you can open on any device. The one you just finished stays on this browser only, so save now if you want the next one to travel.

Not
now

you@example.com

6-digit PIN

Save progress



[Back](#) [Account](#)