

CCAF



Pass

91.67% (55/60)

DOMAIN BREAKDOWN

research_pipeline	14/17 (82%)
extraction_pipeline	16/16 (100%)
code_exploration	13/15 (87%)
customer_support	12/12 (100%)

Share

Open saved result

Missed (5)

All (60)

research_pipeline

Your multi-agent research pipeline crashed after processing 12 of 28 documents. The web search agent had identified relevant sources, the document analysis agent had partially completed extraction, and the synthesizer had begun pattern identification.

You need to resume processing without repeating work or losing fidelity of prior findings.

What state management approach best balances information fidelity with context efficiency when restoring agent state?

Your answer: B. Persist the coordinator's conversation log containing all task delegations and responses, providing this to agents when resuming.

Correct: C. Have each agent persist a structured report to a known location. On resume, the coordinator loads the reports and injects relevant state into agent prompts.

code_exploration

Your agent needs to insert a new helper function into the middle of a 150-line utility module, between two existing functions. The Edit tool fails because its 'old_string' parameter cannot find unique text to match — the file has repetitive docstrings, variable names, and structural patterns.

What's the most reliable way to complete this insertion?

Your answer: B. Use Edit's 'replace_all' parameter to target a common pattern and embed the new function in the replacement text

Correct: D. Use Read to load the file, add the function at the appropriate location, then Write the updated file

research_pipeline

Production monitoring shows that follow-up queries like "summarize what we learned about market trends" consistently take 40+ seconds. Investigation reveals the coordinator spawns the synthesis subagent for each summarization request, passing 80K+ tokens of accumulated findings. The coordinator already has these findings in its context from orchestrating the research.

What's the most effective way to improve response time for these follow-up summaries?

Your answer: C. Enable prompt caching on the synthesis subagent to reduce the overhead of repeatedly transferring the same research findings.

Correct: B. Have the coordinator handle straightforward summarization requests directly using its existing context, reserving subagent spawning for complex analysis.

research_pipeline

In production, final reports frequently contain claims without proper source attribution. Investigation shows that while the web search and document analysis agents correctly attach citations to their outputs, the synthesis agent loses track of which sources support which conclusions when combining findings.

What's the most effective architectural change?

Your answer: D. Have the coordinator inject source identifier prefixes into text before each handoff, then parse these prefixes at report generation to reconstruct citations.

Correct: B. Require all subagents to output structured claim-source mappings that the synthesis agent must preserve and merge when combining findings from multiple sources.

code_exploration

An engineer asks your agent to identify untested code paths in a legacy payment processing module spanning 45 files. After reading the first 8 source files, the agent's responses are becoming noticeably less accurate—it's forgetting previously discussed code patterns and hasn't yet located all test files or traced critical payment flows.

What's the most effective approach to complete this investigation?

Your answer: D. Switch to using Grep to search for specific function names instead of reading full files, reducing the content loaded into context for remaining exploration.

Correct: B. Spawn subagents to investigate specific questions (e.g., "find all test files for payment processing", "trace refund flow dependencies") while the main agent coordinates findings and preserves high-level understanding.



Save your progress from here
Email + a 6-digit PIN attaches **future** sittings to an account you can open on any device. The one you just finished stays on this browser only, so save now if you want the next one to travel.

Not
now

you@example.com

6-digit PIN

Save progress

 Success!


[Privacy](#) · [Help](#)

CyberSkill is an independent practice-exam resource and is neither affiliated with, nor authorized, sponsored, or approved by, Anthropic, PBC.

[Back](#) [Account](#)